

# MH Examples: Logistic Regression

Chang Tu

```
###MH Examples: Logistic Regression
```

```
library(MASS)
library(mvtnorm)      #mutivariate normal density function and random variate generation
library(LearnBayes)   #laplace function
library(arm)          #logit, invlogit, display and bayesglm functions
```

```
## Loading required package: Matrix
```

```
## Loading required package: lme4
```

```
##
## arm (Version 1.12-2, built: 2021-10-15)
```

```
## Working directory is /work/files/workspace/Baysian-inference/CODE
```

```
#####
#####      Logistic regression example #####
#####
```

```
##If not using a Rstudio project, unhash the line below and change to the
##name of the folder you are using.
```

```
#setwd("~/Patrick/Stat314-2018")
```

```
##read in additional functions to help with logistic
##regression
##The code below assumes a folder structure with sub-folders "R" for code
##and "data". If you do not have this set-up you will need to modify
##the lines below appropriately.
```

```
source("~/workspace/assignment4/logisticregression_functions.r")
readdata <- read.csv("~/workspace/assignment4/nzis11-cart-surf.csv",header=TRUE)
str(readdata)
```

```
## 'data.frame':    29471 obs. of  8 variables:
## $ ethnicity      : int  1 1 1 3 1 1 1 1 1 2 ...
## $ lgr            : int  2 2 10 2 2 5 10 10 10 1 ...
## $ sex            : int  2 2 2 2 2 2 1 1 1 2 ...
## $ agegrp         : int  45 65 40 15 30 65 20 50 40 45 ...
## $ qualification: int  3 1 3 2 4 1 2 4 3 3 ...
## $ occupation     : int  2 10 2 10 2 7 3 1 2 4 ...
## $ hours          : num  30 0 15 0 37.5 16 0 45 40 40 ...
## $ income         : int  644 0 2671 458 1019 1033 0 2405 576 540 ...
```

```
##To check functions have been read in:  
logpost_logit_norm #code should appear in the console
```

```
## function (beta, prior_mean, prior_variance, model)  
## {  
##   log_posterior <- loglike_logit(beta = beta, model = model) +  
##     dmvnorm(x = beta, mean = prior_mean, sigma = prior_variance,  
##           log = TRUE)  
##   return(log_posterior)  
## }
```

```
mylaplace #code should appear in the console
```

```
## function (logpost, mode, maxiter = 1000, ...)  
## {  
##   options(warn = -1)  
##   fit = optim(mode, logpost, gr = NULL, ..., hessian = TRUE,  
##             control = list(fnscale = -1, maxit = maxiter))  
##   options(warn = 0)  
##   mode = fit$par  
##   h = -solve(fit$hessian)  
##   p = length(mode)  
##   int = p/2 * log(2 * pi) + 0.5 * log(det(h)) + logpost(mode,  
##     ...)  
##   stuff = list(mode = mode, var = h, int = int, converge = fit$convergence ==  
##     0, code = fit$convergence, counts = fit$counts)  
##   return(stuff)  
## }
```

```
###set-up variables
```

```
attach(readdata)  
highincome <- as.numeric(readdata$income > 1250)  
str(highincome)
```

```
## num [1:29471] 0 0 1 0 0 0 0 1 0 0 ...
```

```
mean(highincome)
```

```
## [1] 0.145974
```

```
agefactor <- relevel(as.factor(readdata$agegrp), "40")  
regfactor <- relevel(as.factor(readdata$lgr), "2") #Auckland as ref  
sexfactor <- relevel(as.factor(readdata$sex), "1") #male as ref  
qualfactor <- relevel(as.factor(readdata$qualification), "2") #school quals as ref  
table(qualfactor)
```

```
## qualfactor  
## 2 1 3 4 5  
## 7064 6891 8435 5223 1858
```

```

occfactor <- relevel(as.factor(readdata$occupation),"1")
table(occfactor)

```

```

## occfactor
##      1      2      3      4      5      6      7      8      9     10
## 3162  4540  2377  1734  2126  1688  1049  2154    24 10617

```

```

obshours_centred <- readdata$hours-mean(readdata$hours)
obshours_cen40 <- readdata$hours-40
hourscut <- cut(readdata$hours,breaks=c(0,10,30,40,50,168),
               include.highest=TRUE,right=FALSE)
table(hourscut)

```

```

## hourscut
##  [0,10)  [10,30)  [30,40)  [40,50)  [50,168)
##   14637    2799    2920    7468    1646

```

```

hoursfactor <- relevel(hourscut,"[0,10)")
table(hoursfactor)

```

```

## hoursfactor
##  [0,10)  [10,30)  [30,40)  [40,50)  [50,168)
##   14637    2799    2920    7468    1646

```

```

### Fit a basic version of your model
#assuming uniform prors for beta parameter. You need to replace
#the ... after ~ with your chosen model for the covariates. Check
# the glm syntax if you need to (?glm). If your outcome variable is something
#Other than highincome then that also needs to be replaced in the call to glm

#
###basic models assuming uniform prors for beta parameter
logitmodel <- glm(highincome ~ hoursfactor + sexfactor +
                  qualfactor,family=binomial(link="logit") )

summary(logitmodel)

```

```
##
## Call:
## glm(formula = highincome ~ hoursfactor + sexfactor + qualfactor,
##      family = binomial(link = "logit"))
##
## Deviance Residuals:
##      Min        1Q    Median        3Q        Max
## -1.4466  -0.5945  -0.3321  -0.2685   2.6118
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.02244    0.05531  -54.642 < 2e-16 ***
## hoursfactor[10,30)  0.50728    0.07794   6.508 7.61e-11 ***
## hoursfactor[30,40)  1.57825    0.05934  26.595 < 2e-16 ***
## hoursfactor[40,50)  1.66099    0.04661  35.633 < 2e-16 ***
## hoursfactor[50,168) 2.66743    0.06307  42.294 < 2e-16 ***
## sexfactor2       -0.35466    0.03734  -9.499 < 2e-16 ***
## qualfactor1        0.07248    0.05959   1.216  0.224
## qualfactor3        0.32169    0.05271   6.103 1.04e-09 ***
## qualfactor4        0.96864    0.05390  17.971 < 2e-16 ***
## qualfactor5        0.42156    0.07983   5.281 1.29e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 24500  on 29469  degrees of freedom
## Residual deviance: 20900  on 29460  degrees of freedom
## (1 observation deleted due to missingness)
## AIC: 20920
##
## Number of Fisher Scoring iterations: 5
```

```
display(logitmodel)
```

```
## glm(formula = highincome ~ hoursfactor + sexfactor + qualfactor,
##      family = binomial(link = "logit"))
##              coef.est coef.se
## (Intercept)    -3.02    0.06
## hoursfactor[10,30)  0.51    0.08
## hoursfactor[30,40)  1.58    0.06
## hoursfactor[40,50)  1.66    0.05
## hoursfactor[50,168) 2.67    0.06
## sexfactor2       -0.35    0.04
## qualfactor1        0.07    0.06
## qualfactor3        0.32    0.05
## qualfactor4        0.97    0.05
## qualfactor5        0.42    0.08
## ---
##      n = 29470, k = 10
##      residual deviance = 20899.7, null deviance = 24499.6 (difference = 3599.9)
```

```
##see what the covariance matrix looks like
vcov(logitmodel)
```

```
##          (Intercept) hoursfactor[10,30) hoursfactor[30,40)
## (Intercept)      0.003059595      -1.229195e-03      -1.249327e-03
## hoursfactor[10,30) -0.001229195      6.075317e-03      1.479753e-03
## hoursfactor[30,40) -0.001249327      1.479753e-03      3.521736e-03
## hoursfactor[40,50) -0.001428943      1.388649e-03      1.390640e-03
## hoursfactor[50,168) -0.001540927      1.349618e-03      1.355418e-03
## sexfactor2      -0.000648049      -2.840374e-04      -2.492331e-04
## qualfactor1      -0.001822761      1.098990e-04      1.055694e-04
## qualfactor3      -0.001695912      -7.916709e-05      -6.247676e-05
## qualfactor4      -0.001611610      -1.811269e-04      -1.715305e-04
## qualfactor5      -0.001714780      -1.687519e-05      -3.747820e-05
##          hoursfactor[40,50) hoursfactor[50,168)      sexfactor2
## (Intercept)      -1.428943e-03      -1.540927e-03 -6.480490e-04
## hoursfactor[10,30)      1.388649e-03      1.349618e-03 -2.840374e-04
## hoursfactor[30,40)      1.390640e-03      1.355418e-03 -2.492331e-04
## hoursfactor[40,50)      2.172871e-03      1.437798e-03  1.615546e-04
## hoursfactor[50,168)      1.437798e-03      3.977579e-03  2.687971e-04
## sexfactor2      1.615546e-04      2.687971e-04  1.394000e-03
## qualfactor1      9.675295e-05      1.193966e-04  3.058381e-05
## qualfactor3      -1.111677e-04      -3.141968e-05  6.280221e-05
## qualfactor4      -1.386672e-04      9.785551e-06 -6.896867e-05
## qualfactor5      -4.346861e-05      -1.087340e-06  1.649586e-05
##          qualfactor1      qualfactor3      qualfactor4      qualfactor5
## (Intercept)      -1.822761e-03 -1.695912e-03 -1.611610e-03 -1.714780e-03
## hoursfactor[10,30)      1.098990e-04 -7.916709e-05 -1.811269e-04 -1.687519e-05
## hoursfactor[30,40)      1.055694e-04 -6.247676e-05 -1.715305e-04 -3.747820e-05
## hoursfactor[40,50)      9.675295e-05 -1.111677e-04 -1.386672e-04 -4.346861e-05
## hoursfactor[50,168)      1.193966e-04 -3.141968e-05  9.785551e-06 -1.087340e-06
## sexfactor2      3.058381e-05  6.280221e-05 -6.896867e-05  1.649586e-05
## qualfactor1      3.550481e-03  1.728839e-03  1.723570e-03  1.730422e-03
## qualfactor3      1.728839e-03  2.778333e-03  1.735850e-03  1.735070e-03
## qualfactor4      1.723570e-03  1.735850e-03  2.905267e-03  1.734567e-03
## qualfactor5      1.730422e-03  1.735070e-03  1.734567e-03  6.373007e-03
```

```
bayesmodell1 <- bayesglm(highincome ~ hoursfactor + sexfactor +
                        qualfactor,family=binomial(link="logit") )

display(bayesmodell1)
```

```
## bayesglm(formula = highincome ~ hoursfactor + sexfactor + qualfactor,
##          family = binomial(link = "logit"))
##          coef.est coef.se
## (Intercept)      -3.02    0.06
## hoursfactor[10,30)   0.50    0.08
## hoursfactor[30,40)   1.58    0.06
## hoursfactor[40,50)   1.66    0.05
## hoursfactor[50,168)  2.66    0.06
## sexfactor2      -0.35    0.04
## qualfactor1         0.07    0.06
## qualfactor3         0.32    0.05
## qualfactor4         0.97    0.05
## qualfactor5         0.42    0.08
## ---
## n = 29470, k = 10
## residual deviance = 20899.7, null deviance = 24499.6 (difference = 3599.9)
```

```
##re-fit controlling for occupation
logitmodel2 <- glm(highincome ~ hoursfactor + sexfactor +
                    qualfactor + occfactor,
                    family=binomial(link="logit") )

#summary(logitmodel2)
display(logitmodel2)
```

```
## glm(formula = highincome ~ hoursfactor + sexfactor + qualfactor +
##      occfactor, family = binomial(link = "logit"))
##              coef.est coef.se
## (Intercept)    -1.31    0.07
## hoursfactor[10,30) -0.46    0.08
## hoursfactor[30,40)  0.62    0.07
## hoursfactor[40,50)  0.71    0.05
## hoursfactor[50,168) 1.72    0.07
## sexfactor2      -0.29    0.04
## qualfactor1      0.16    0.06
## qualfactor3      0.20    0.05
## qualfactor4      0.44    0.06
## qualfactor5      0.31    0.08
## occfactor2      -0.11    0.05
## occfactor3      -0.87    0.07
## occfactor4      -1.30    0.09
## occfactor5      -0.93    0.08
## occfactor6      -1.01    0.08
## occfactor7      -0.99    0.09
## occfactor8      -1.52    0.08
## occfactor9      -1.48    0.63
## occfactor10     -2.96    0.10
## ---
##      n = 29470, k = 19
##      residual deviance = 19296.4, null deviance = 24499.6 (difference = 5203.2)
```

```
###illustrate how glm deals with factors
##extract model matrix -- internal glm representation of the model structure
chkX <- model.matrix(logitmodel)
head(chkX)
```

```
##      (Intercept) hoursfactor[10,30) hoursfactor[30,40) hoursfactor[40,50)
## 1             1             0             1             0
## 2             1             0             0             0
## 3             1             1             0             0
## 4             1             0             0             0
## 5             1             0             1             0
## 6             1             1             0             0
##      hoursfactor[50,168) sexfactor2 qualfactor1 qualfactor3 qualfactor4
## 1             0             1             0             1             0
## 2             0             1             1             0             0
## 3             0             1             0             1             0
## 4             0             1             0             0             0
## 5             0             1             0             0             1
## 6             0             1             1             0             0
##      qualfactor5
## 1             0
## 2             0
## 3             0
## 4             0
## 5             0
## 6             0
```

```
##Should really control for age also; income changes with age and
#age is also #related to quals
```

```
##re-fit controlling for occupation and age
logitmodel3 <- glm(highincome ~ hoursfactor + sexfactor +
                    qualfactor + occfactor + agefactor,
                    family=binomial(link="logit") )

summary(logitmodel3)
```

```
##
## Call:
## glm(formula = highincome ~ hoursfactor + sexfactor + qualfactor +
##      occfactor + agefactor, family = binomial(link = "logit"))
##
## Deviance Residuals:
##      Min        1Q      Median        3Q        Max
## -1.7130   -0.5904   -0.2187   -0.1199    3.4078
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -0.891298    0.084653  -10.529 < 2e-16 ***
## hoursfactor[10,30) -0.370929    0.083712   -4.431 9.38e-06 ***
## hoursfactor[30,40)  0.718255    0.066755   10.760 < 2e-16 ***
## hoursfactor[40,50)  0.789141    0.053998   14.614 < 2e-16 ***
## hoursfactor[50,168) 1.813919    0.071013   25.543 < 2e-16 ***
## sexfactor2      -0.317168    0.040876   -7.759 8.54e-15 ***
## qualfactor1     -0.006148    0.064174   -0.096  0.92367
## qualfactor3     -0.007392    0.056604   -0.131  0.89610
## qualfactor4      0.282402    0.061270    4.609 4.04e-06 ***
## qualfactor5      0.084274    0.084499    0.997  0.31860
## occfactor2      -0.090545    0.055106   -1.643  0.10036
## occfactor3      -0.854089    0.068823  -12.410 < 2e-16 ***
## occfactor4      -1.315746    0.087537  -15.031 < 2e-16 ***
## occfactor5      -0.978108    0.076708  -12.751 < 2e-16 ***
## occfactor6      -0.992900    0.084085  -11.808 < 2e-16 ***
## occfactor7      -1.010010    0.093423  -10.811 < 2e-16 ***
## occfactor8      -1.521534    0.085242  -17.850 < 2e-16 ***
## occfactor9      -1.418897    0.636817   -2.228  0.02587 *
## occfactor10     -2.824018    0.103850  -27.193 < 2e-16 ***
## agefactor15     -1.764882    0.174267  -10.127 < 2e-16 ***
## agefactor20     -1.407822    0.097738  -14.404 < 2e-16 ***
## agefactor25     -0.898490    0.083843  -10.716 < 2e-16 ***
## agefactor30     -0.463247    0.078323   -5.915 3.33e-09 ***
## agefactor35     -0.011822    0.074081   -0.160  0.87321
## agefactor45     -0.095715    0.073016   -1.311  0.18990
## agefactor50     -0.034945    0.073904   -0.473  0.63632
## agefactor55     -0.225836    0.079658   -2.835  0.00458 **
## agefactor60     -0.170303    0.082538   -2.063  0.03908 *
## agefactor65     -0.122871    0.091067   -1.349  0.17726
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 24500  on 29469  degrees of freedom
## Residual deviance: 18745  on 29441  degrees of freedom
## (1 observation deleted due to missingness)
## AIC: 18803
##
## Number of Fisher Scoring iterations: 7
```

```
display(logitmodel3)
```



```
## glm(formula = highincome ~ hoursfactor + sexfactor + qualfactor +
##      occfactor + agefactor, family = binomial(link = "logit"))
##              coef.est coef.se
## (Intercept)    -0.89    0.08
## hoursfactor[10,30) -0.37    0.08
## hoursfactor[30,40)  0.72    0.07
## hoursfactor[40,50)  0.79    0.05
## hoursfactor[50,168) 1.81    0.07
## sexfactor2       -0.32    0.04
## qualfactor1      -0.01    0.06
## qualfactor3      -0.01    0.06
## qualfactor4       0.28    0.06
## qualfactor5       0.08    0.08
## occfactor2       -0.09    0.06
## occfactor3       -0.85    0.07
## occfactor4      -1.32    0.09
## occfactor5      -0.98    0.08
## occfactor6      -0.99    0.08
## occfactor7      -1.01    0.09
## occfactor8      -1.52    0.09
## occfactor9      -1.42    0.64
## occfactor10     -2.82    0.10
## agefactor15     -1.76    0.17
## agefactor20     -1.41    0.10
## agefactor25     -0.90    0.08
## agefactor30     -0.46    0.08
## agefactor35     -0.01    0.07
## agefactor45     -0.10    0.07
## agefactor50     -0.03    0.07
## agefactor55     -0.23    0.08
## agefactor60     -0.17    0.08
## agefactor65     -0.12    0.09
## ---
##      n = 29470, k = 29
##      residual deviance = 18745.1, null deviance = 24499.6 (difference = 5754.5)
```

```
####Specify priors:
#####
##Analysis with an informative prior
#####

##obtain priors for parameters of logitmodel
##using the makepriorfunctions makepriorb0() and makepriorbreg() read in from
#logisticregression_functions.r

##Intercept - specify prior interval on the probability scale
b0prior <- makepriorb0(low=0.01,high=0.10,credibility=0.95)
b0prior
```

```
## [1] -3.3961722  0.6117192
```

```
##logistic regression parameters
##specify prior on odds ratio scale; makepriorreg converts to prior
#on the beta scale
bfemale    <- makepriorbreg(low=0.5,high=1.1,credibility=0.95)
bqualnone  <- makepriorbreg(low=0.4,high=2,credibility=0.95)
bqualtrade <- makepriorbreg(low=0.7,high=3,credibility=0.95)
bqualuni   <- makepriorbreg(low=0.7,high=4,credibility=0.95)
bqualother <- makepriorbreg(low=0.1,high=10,credibility=0.95)
bhours10_30 <- makepriorbreg(low=1,high=5,credibility=0.9)
bhours30_40 <- makepriorbreg(low=1,high=5,credibility=0.9)
bhours40_50 <- makepriorbreg(low=1,high=5,credibility=0.9)
bhours50pl <- makepriorbreg(low=1.1,high=10,credibility=0.9)

###add lines to set priors for any additional parameters
###check a few of these
bfemale
```

```
## [1] -0.2989185  0.2011408
```

```
bqualuni
```

```
## [1] 0.5148097 0.4446432
```

```
##combine prior means into a vector and construct a
#diagonal prior variance matrix
##based on the prior standard deviations prior.matrix

#First put everything in matrix - column1= priormean,
#column 2 = prior sd
prior.matrix<- rbind(b0prior,bhours10_30,bhours30_40,bhours40_50,
                    bhours50pl,bfemale,bqualnone,bqualtrade,
                    bqualuni,bqualother)

prior.matrix
```

```
##           [,1]      [,2]
## b0prior    -3.396172e+00 0.6117192
## bhours10_30 8.047190e-01 0.4892344
## bhours30_40 8.047190e-01 0.4892344
## bhours40_50 8.047190e-01 0.4892344
## bhours50pl  1.198948e+00 0.6709639
## bfemale    -2.989185e-01 0.2011408
## bqualnone   -1.115718e-01 0.4105784
## bqualtrade  3.709687e-01 0.3712536
## bqualuni    5.148097e-01 0.4446432
## bqualother  2.220446e-16 1.1748099
```

```
prior_mean_inf <- prior.matrix[,1]    #prior mean vector
prior_variance_inf <- diag((prior.matrix[,2])^2) #prior variance matrix

prior_mean_inf
```

```
##          b0prior    bhours10_30    bhours30_40    bhours40_50    bhours50pl
## -3.396172e+00  8.047190e-01  8.047190e-01  8.047190e-01  1.198948e+00
##          bfemale    bqalnone    bqaltrade    bqaluni    bqalother
## -2.989185e-01 -1.115718e-01  3.709687e-01  5.148097e-01  2.220446e-16
```

```
prior_variance_inf
```

```
##          [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.3742004 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## [2,] 0.0000000 0.2393503 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## [3,] 0.0000000 0.0000000 0.2393503 0.0000000 0.0000000 0.0000000 0.0000000
## [4,] 0.0000000 0.0000000 0.0000000 0.2393503 0.0000000 0.0000000 0.0000000
## [5,] 0.0000000 0.0000000 0.0000000 0.0000000 0.4501926 0.0000000 0.0000000
## [6,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.04045761 0.0000000
## [7,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.1685747
## [8,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## [9,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
## [10,] 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
##          [,8]      [,9]      [,10]
## [1,] 0.0000000 0.0000000 0.0000000
## [2,] 0.0000000 0.0000000 0.0000000
## [3,] 0.0000000 0.0000000 0.0000000
## [4,] 0.0000000 0.0000000 0.0000000
## [5,] 0.0000000 0.0000000 0.0000000
## [6,] 0.0000000 0.0000000 0.0000000
## [7,] 0.0000000 0.0000000 0.0000000
## [8,] 0.1378292 0.0000000 0.0000000
## [9,] 0.0000000 0.1977076 0.0000000
## [10,] 0.0000000 0.0000000 1.380178
```

```
#####
#Next step is build an approximation to the posterior
#which we can use to generate starting values for the Markov Chains
#We can also base the variance of the jumping density on an
#approximation to the posterior variance
#We will use a multivariate normal approximation centred on the
#posterior model with variance determined by the curvature of the
#log-posterior at the mode. The posterior mode and approximate variance
#can be obtained using either the laplace() function or
#bayesglm(), given appropriate specification of the prior in both cases )

##test out the log posterior function
bmle <- coef(logitmodel) #For convenience evaluate the unnormalised
#log posterior at the mle of the basic model
#called logitmodel
testpost <- logpost_logit_norm(beta=bmle,
                                prior_mean=prior_mean_inf,
                                prior_variance=prior_variance_inf,
                                model=logitmodel)

testpost
```

```
## [1] -10458.15
```

```
##use laplace to find the posterior mode and a variance approximation
## first argument passed to laplace is the name of the log-posterior function,
#then an initial guess at the mode --suggest using the mle from the glm);
#then all the arguments required by the log posterior function need to be
#explicitly passed to laplace()

laplace_post <- laplace(logpost_logit_norm,mode=bmle,
                        prior_mean=prior_mean_inf,
                        prior_variance=prior_variance_inf,
                        model=logitmodel)
##in the above call to laplace() the parameter mode just specified
##the starting values for the optimisation procedure; ideally an
##informed guessfor the posterior mode would be supplied
laplace_post$converge      ##This should equal TRUE otherwise
```

```
## [1] TRUE
```

```
##there is a problem. For a complex model the
##solution may simply be to increase the number
##of iterations in laplace. This can achieved
##using the mylaplace() function in place of
##laplace.This has the same syntax, except
##for an additional
# parameter, maxiter. The default maximum number of iterations
#in mylaplace is set to
#1000, in contrast to laplace in which the maximum
#number of iterations is set to 500.
```

```
mean_approx <- laplace_post$mode
var_approx <- laplace_post$var
sd_approx <- sqrt(diag(var_approx))
###display mode and approximation to variance
cbind(mean_approx,sd_approx)
```

```
##          mean_approx  sd_approx
## (Intercept)    -3.00997705 0.05412000
## hoursfactor[10,30)  0.50787063 0.07655976
## hoursfactor[30,40)  1.56615282 0.05865690
## hoursfactor[40,50)  1.65350144 0.04603601
## hoursfactor[50,168) 2.65393987 0.06248255
## sexfactor2       -0.35184188 0.03665130
## qualfactor1       0.06360094 0.05850864
## qualfactor3       0.31373853 0.05171067
## qualfactor4       0.96093295 0.05297908
## qualfactor5       0.41351140 0.07919749
```



```

##evaluate unnormalized posterior at current value

logpost_old <- logpost_logit_norm(beta=oldbeta,
                                prior_mean=prior_mean_inf,
                                prior_variance=prior_variance_inf,
                                model=logitmodel)

##compute MH acceptance ratio or log MH acceptance ratio

logrMH <- logpost_prop - logpost_old

#decide whether to accept the proposal: create an acceptance indicator
#called accept
logU <- log(runif(1))
accept <- (logU < logrMH)

#Assuming acceptance ratio is called accept
storeaccept[i,j] <- as.numeric(accept)

###If jump accepted update oldbeta to newbeta
if (accept) {
  oldbeta <- newbeta
}
##else just stay at curent value of oldbeta
} ###end loop over simulations

} ###end loop over chains
) #end(system.time)

```

```

##      user  system elapsed
## 26.504    0.204   26.713

```

```

#Rhat statistics

##Simply loop over the parameters and carry out the computations
# as per Gelman et al

#decide on burn-in period

burnin <- 250
npos <- nsim - burnin  ##size of retained posterior sample in each chain

#set-up a vector for storing the Rhat statistics
rhat <- vector(mode="numeric",length=length(mean_approx))

#set-up a vector for storing the effective sample sizes
neff <- vector(mode="numeric",length=length(mean_approx))

for (k in 1:length(bmle)) {
  ##subsets out values for ith parameter

  betak <- storebeta[k,1:nchains,((burnin+1):nsim)]  ##subsets out post burnin sample
  for kth parameter
  chainmeans <- rowMeans(betak)
  withinstd <- apply(betak,MARGIN=1,FUN="sd")
  betweenstd <- sd(chainmeans)
  B = npos*betweenstd^2
  W = mean(withinstd^2)
  varplus <- ((npos-1)/npos) * W + (1/npos)*B
  rhat[k] <- sqrt(varplus/W)
  neff[k] <- nchains*npos*(varplus/B)
}
rhat

```

```

## [1] 1.0073349 0.9996836 1.0150601 1.0020650 1.0208741 1.0236054 1.0279172
## [8] 1.0712192 1.0268253 1.0434332

```

```

neff  ##Some of these are concerningly low - we would probably want to

```

```

## [1] 193.60475 8163.18265 98.60730 586.72998 72.40098 64.45433
## [7] 55.01874 23.18039 57.12667 36.39167

```


```
##increase the posterior sample size in reality

##Traceplots
##useful but optional for assignment    use code like this
##Assumes 3 chains and looking at just the 2nd parameter
##Really needs to be repeated for each parameter.

##You may also want to exclude the burnin
x <- seq(from=1,to=nsim,by=1)

plot(x,storebeta[2,1,],col="black",type="l",
      main="Traceplot for hours10_30")
lines(x,storebeta[2,2,],col="red")
lines(x,storebeta[2,3,],col="blue")
```

**Traceplot for hours10\_30**





```

####If satisfied about convergence use the post burn-in sample for posterior inference

possamp <- storebeta[,,(burnin+1):nsim] # this is our sample from the posterior

##produce posterior density plot for at least one of the logistic regression coefficients
##credible interval; posterior probability > 0.
##To subset out the posterior for a single
##parameters use code like pos_betak <- possamp[k,,];
#Functions like summary() and quantile() can then be applied
##to pos_betak and work happily even though it is an
#array rather than a vector

##Up to the user to know which parameters are of particular interest
pos_tertiary <- possamp[9,,]


mean(pos_tertiary > 0)

```

```
## [1] 1
```

```
plot(density(pos_tertiary),main="posterior density for tertiary ed log-odds ratio ")
```

### posterior density for tertiary ed log-odds ratio




```

##back on the odds ratio scale
plot(density(exp(pos_tertiary),adjust=1.7),
     main="posterior density for tertiary ed odds ratio ",
     xlab="odds ratio")

```

## posterior density for tertiary ed odds ratio



```
pos_tert <- possamp[9,,]  
#####90% credible interval for tertiary education odds ratio  
quantile(exp(pos_tertiary),probs=c(0.05,0.5,0.95))
```

```
##          5%          50%          95%  
## 2.433730 2.631385 2.883247
```

```
##clearly seems to be a fairly strong tertiary education effect - but recall this  
##is without adjusting for age or occupation.
```